

# COMP348 — Document Processing and the Semantic Web

## Week 05 Lecture 1: Processing Text Sequences

Diego Mollá

Department of Computer Science  
Macquarie University

COMP348 2019H1

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

## Reading

- Deep Learning book, chapter 6.
- Understanding LSTM Networks,  
<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>.

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

# Words as Arbitrary Symbols

- Words are encoded as arbitrary symbols.
- Within one language there is no clear correspondence between a word symbol and its meaning.
  - “dig” vs. “dog”
  - “car” vs. “automobile”
- Different languages may use different representations of the same word.



[https://en.wikipedia.org/wiki/File:Hello\\_in\\_different\\_languages\\_](https://en.wikipedia.org/wiki/File:Hello_in_different_languages_)

# Ambiguities Everywhere

Language features ambiguity at multiple levels.


## Lexical Ambiguity

Example from Google's dictionary:

- bank (n): the land alongside or sloping down a river or lake.
- bank (n): financial establishment that uses money deposited by customers for investment, . . .
- bank (v): form in to a mass or mound.
- bank (v): build (a road, railway, or sports track) higher at the outer edge of a bend to facilitate fast cornering.
- . . .

# So many words!

- Any language features a large number of distinct words.
- New words are coined.
- Words change their use in time.
- There are also names, numbers, dates... an infinite number.



<https://trends.google.com>

# Long-distance Dependencies

- Sentences are sequences of words.
- Words close in the sentence are often related.
- But sometimes there are relations between words far apart.

grammatical: “The man living upstairs is very cheerful”

“The people living upstairs are very cheerful”

reference: “I bought a book from the bookshp and I liked it”

knowledge: “I was born in France and I speak fluent French”



# Programme

## 1 Word Embeddings


- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences


- Modelling Text Sequences
- Sequence Labelling

# Word Embeddings


- First introduced in 2013, nowadays is one of the most common ingredients in text processing systems.
- Word embeddings squarely aim at addressing the issue of representing words as continuous vectors of integers.
- Words with similar context are mapped to similar vectors.
- Embeddings are learnt using large, unlabelled training data.



Male-Female



Verb tense




Country-Capital

<https://www.tensorflow.org/tutorials/representation/word2vec>

# One-hot vs. word embeddings


## One-hot

- Sparse
- Binary values (typically)
- High-dimensional
- Hard-coded



## Word embeddings

- Dense
- Continuous values
- Lower-dimensional
- Learned from data



# Two Ways to Obtain Word Embeddings

- 1 Learn the word embeddings jointly with the task you care about (e.g. document classification).
- 2 Use pre-trained word embeddings.

# Learning Word Embeddings

- You can add a dense layer as the first layer of your network and let the system learn the optimal weights.
- This approach is so useful and common that many deep learning frameworks define an “embedding” layer that facilitates this.
- The input to the “embedding” layer is the word index.
- The output is the word embedding.

# Using pre-trained word embeddings

- Sometimes we have so little training data that many words are poorly represented.
- Often, words in the training data do not occur in the test data.
- For these words we would not be able to learn the embeddings.
- Several people have computed word embeddings for large vocabularies using large data sets.
- We can then use these pre-trained embeddings to map from the word index to the word embedding.

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings


## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling



# Handling Text Sequences


- A document is a sequence of words.
- Many document representations are based on a bag-of-words approach.
  - Word order is ignored.
- Even word embeddings ignore word order.
- A **Recurrent Neural Network** (RNN), however, is designed to process sequences.



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# A Recurrent Neural Network


- A RNN is a neural network that is composed of RNN cells.
- Each RNN cell takes as input two pieces of information:
  - 1 A vector representing an item in the sequence.
  - 2 The state resulting from processing the previous items.
- The output of the RNN cell is a state that can be fed to the next cell in the sequence.
- All RNN cells are identical copies. In a sense, we can say that an RNN cell is the same for all words in the sequence.



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# Recurrent Neural Networks

- RNNs are designed to model long-distance dependencies.
- Vanilla RNN cells were used since 1990s.




<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

# LSTMs and GRUs

- Vanilla RNN cells are still too simple and they do not memorise long-distance dependencies easily.
- More complex RNN cells have been designed specifically to address this issue.
- These RNN cells have components that are programmed to memorise or forget past information.
- Current most popular RNN cells are:

**LSTM** Long Short Term Memory (picture).

**GRU** Gated Recurrent Unit; a more recent, simpler cell.



# RNNs in Practice

- Most deep learning frameworks include special layers for RNNs.
- When you use an RNN layer, you have the option to specify the type of RNN cell.
- You often have the option to use the state of the last cell, or the state of all cells.

# Programme

## 1 Word Embeddings

- Challenges of Text for Machine Learning
- Word Embeddings

## 2 Text Sequences

- Modelling Text Sequences
- Sequence Labelling

# What is Sequence Labelling?

- A **sequence labelling** problem is one where:
  - the input consists of a sequence  $\mathbf{X} = (X_1, \dots, X_n)$ , and
  - the output consists of a sequence  $\mathbf{Y} = (Y_1, \dots, Y_n)$  of labels, where:
    - $Y_i$  is the label for element  $X_i$
- Example: Part-of-speech tagging

$$\begin{pmatrix} \mathbf{Y} \\ \mathbf{X} \end{pmatrix} = \begin{pmatrix} \text{Verb,} & \text{Determiner,} & \text{Noun} \\ \text{spread,} & \text{the,} & \text{butter} \end{pmatrix}$$

- Example: Spelling correction

$$\begin{pmatrix} \mathbf{Y} \\ \mathbf{X} \end{pmatrix} = \begin{pmatrix} \text{write,} & \text{a,} & \text{book} \\ \text{rite,} & \text{a,} & \text{buk} \end{pmatrix}$$

# Other applications of sequence labelling

- **Named entity recognition** and classification (NER) involves finding the named entities in a text and identifying what type of entity they are (e.g., person, location, corporation, dates, etc.).
- **Speech transcription** can be seen as a sequence labelling task:
  - The input  $\mathbf{X} = (X_1, \dots, X_n)$  is a sequence of **acoustic frames**  $X_i$ , where  $X_i$  is a set of features extracted from a 50msec window of the speech signal.
  - The output  $\mathbf{Y}$  is a sequence of words (the transcript of the speech signal).
- **Financial applications** of sequence labelling:
  - Identifying trends in price movements.
- **Biological applications** of sequence labelling:
  - Gene-finding in DNA or RNA sequences.



# Sequence Labelling as Classification I

## Can we just use a standard classifier?

- Standard classifiers (K-Nearest Neighbours, Naïve Bayes, Support Vector Machine, ...) assume **independence between samples**:
  - The probability of the label assigned to sample  $i$  is independent to the probability of the label assigned to sample  $j$ .
- But in sequence labelling there is interdependence between the labels of different samples.

# Modelling Context

## Classifier with context features

- A (crude) approach to model interdependence between samples is to add context features.
- For example, we can use features based on previous words and following words.
- We can even incorporate the label of the previous word as a feature.
- But it is not so easy to incorporate the label of **both** the previous word and the following word.

# Using Recurrent Neural Networks for Sequence Labelling I

- We have seen how RNNs can be used to classify documents.
- Similarly, we can use RNNs to classify sequences of words.
- In Keras, we can define a dense layer that is repeated at the output of each recurrent cell.

```
model = Sequential()  
model.add(Embedding(max_words, 32))  
model.add(LSTM(32, return_sequences=True))  
model.add(TimeDistributed(Dense(num_tags,  
                                activation='softmax')))  
  
model.compile(optimizer='Adam',  
              loss='categorical_crossentropy',  
              metrics=['acc'])
```

# Using Recurrent Neural Networks for Sequence Labelling II

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, None, 32)	320000
lstm_1 (LSTM)	(None, None, 32)	8320
time_distributed_1 (TimeDist	(None, None, 20)	660

# Take-home Messages

- 1 Explain some of the fundamental challenges that plain text represents to machine learning.
- 2 Apply word embeddings in deep learning.
- 3 Use recurrent neural networks for text classification.
- 4 Comment on the issues of sequence labelling.

# What's Next

## Week 6

- Generating text.
- Reading: Deep Learning book, chapter 8.1